

Simple Housing Dataset

Create a regression model that predicts the price of boston house

```
!wget https://storage.googleapis.com/nicksdemobucket/housing-data.csv
```

```
--2019-03-30 11:55:25-- https://storage.googleapis.com/nicksdemobucket/housing-data.csv
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.141.128, 2607:f8t
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.141.128|:443...
HTTP request sent, awaiting response... 200 OK
Length: 816 [application/octet-stream]
Saving to: 'housing-data.csv'
```

```
housing-data.csv 100%[=====>] 816 --.-KB/s in 0s
```

```
2019-03-30 11:55:25 (15.1 MB/s) - 'housing-data.csv' saved [816/816]
```

```
import pandas as pd
df= pd.read_csv('housing-data.csv')
```

```
df.head()
```

```
sqft  bdrms  age  price
0  2104    3   70  399900
1  1600    3   28  329900
2  2400    3   44  369000
3  1416    2   49  232000
4  3000    4   75  539900
```

```
df.describe()
```

	sqft	bdrms	age	price
count	47.000000	47.000000	47.000000	47.000000
mean	2000.000000	2.170000	40.714000	240140.000000
std	1415.000000	1.000000	18.000000	100000.000000
min	0.000000	0.000000	0.000000	0.000000
max	1415.000000	4.000000	60.000000	1000000.000000

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 47 entries, 0 to 46
Data columns (total 4 columns):
sqft      47 non-null int64
bdrms     47 non-null int64
age       47 non-null int64
price     47 non-null int64
dtypes: int64(4)
memory usage: 1.5 KB
```

```
from sklearn.preprocessing import OneHotEncoder
X = df.iloc[:, :3] #X
X.shape
```

```
(47, 3)
```

```
y = df.iloc[:, 3:] #y
y.shape
```

```
(47, 1)
```

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.train import AdamOptimizer
```

```
model = Sequential()
```

```
model.add(Dense(60, input_shape=(3,), activation='relu'))
model.add(Dense(60, activation='relu'))
model.add(Dense(1))
```

```
model.compile(optimizer=AdamOptimizer(0.01),
              loss='mean_squared_error',
              metrics=['mean_absolute_error'])
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/keras/layers/core.py:107:
Instructions for updating:
Use tf.cast instead.
```

```
h = model.fit(X, y, epochs=100, batch_size=16, validation_split=0.2)
```

Train on 37 samples, validate on 10 samples

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/Instructions for updating:

Use tf.cast instead.

Epoch 1/100

37/37 [=====] - 0s 9ms/sample - loss: 135910616313.0811 - me

Epoch 2/100

37/37 [=====] - 0s 364us/sample - loss: 134570994494.2703 -

Epoch 3/100

37/37 [=====] - 0s 618us/sample - loss: 132700036455.7838 -

Epoch 4/100

37/37 [=====] - 0s 377us/sample - loss: 129851645509.1892 -

Epoch 5/100

37/37 [=====] - 0s 313us/sample - loss: 126279999709.4054 -

Epoch 6/100

37/37 [=====] - 0s 405us/sample - loss: 120869419727.5676 -

Epoch 7/100

37/37 [=====] - 0s 408us/sample - loss: 114874362132.7568 -

Epoch 8/100

37/37 [=====] - 0s 359us/sample - loss: 106137831368.6487 -

Epoch 9/100

37/37 [=====] - 0s 327us/sample - loss: 96033242028.9730 - n

Epoch 10/100

37/37 [=====] - 0s 366us/sample - loss: 83675766894.7027 - n

Epoch 11/100

37/37 [=====] - 0s 328us/sample - loss: 70092994615.3513 - n

Epoch 12/100

37/37 [=====] - 0s 396us/sample - loss: 54862042084.3243 - n

Epoch 13/100

37/37 [=====] - 0s 336us/sample - loss: 39121233975.3514 - n

Epoch 14/100

37/37 [=====] - 0s 298us/sample - loss: 24855652490.3784 - n

Epoch 15/100

37/37 [=====] - 0s 270us/sample - loss: 12803896873.5135 - n

Epoch 16/100

37/37 [=====] - 0s 475us/sample - loss: 5283406612.7568 - me

Epoch 17/100

37/37 [=====] - 0s 382us/sample - loss: 4480017484.1081 - me

Epoch 18/100

37/37 [=====] - 0s 372us/sample - loss: 6552781270.4865 - me

Epoch 19/100

37/37 [=====] - 0s 403us/sample - loss: 9220114653.4054 - me

Epoch 20/100

37/37 [=====] - 0s 465us/sample - loss: 9153664249.0811 - me

Epoch 21/100

37/37 [=====] - 0s 512us/sample - loss: 6952015325.4054 - me

Epoch 22/100

37/37 [=====] - 0s 357us/sample - loss: 4707606458.8108 - me

Epoch 23/100

37/37 [=====] - 0s 518us/sample - loss: 4183094161.2973 - me

Epoch 24/100

37/37 [=====] - 0s 450us/sample - loss: 4604909014.4865 - me

Epoch 25/100

37/37 [=====] - 0s 493us/sample - loss: 4897991908.3243 - me

Epoch 26/100

37/37 [=====] - 0s 400us/sample - loss: 4975622974.2703 - me


Epoch 27/100

```
h.history.keys()
```

```
dict_keys(['loss', 'mean_absolute_error', 'val_loss', 'val_mean_absolute_error'])
```

```
import matplotlib.pyplot as plt
plt.plot(h.history['loss'])
plt.plot(h.history['val_loss'])
```

```
[<matplotlib.lines.Line2D at 0x7f7caea8eba8>]
```



Binary Classification

Create a binary classifier for the titanic dataset, will person x survive?

```
!wget https://storage.googleapis.com/nicksdemobucket/titanic-train.csv
```

```
--2019-03-31 11:38:38-- https://storage.googleapis.com/nicksdemobucket/titanic-train.csv
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.195.128, 2607:f8t
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.195.128|:443...
HTTP request sent, awaiting response... 200 OK
Length: 61194 (60K) [application/octet-stream]
Saving to: 'titanic-train.csv'
```

```
titanic-train.csv 100%[=====>] 59.76K --.-KB/s in 0.001s
```

```
2019-03-31 11:38:38 (93.3 MB/s) - 'titanic-train.csv' saved [61194/61194]
```

```
import pandas as pd
df= pd.read_csv('titanic-train.csv')
```

```
df.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171 7.
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599 71.
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282 51.

```
df.describe()
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	F
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329

```
data = df.drop(['Name', 'Ticket', 'Fare', 'Cabin', 'Embarked', 'Age'],axis =1)
data.head()
```

	PassengerId	Survived	Pclass	Sex	SibSp	Parch
0	1	0	3	male	1	0
1	2	1	1	female	1	0
2	3	1	3	female	0	0
3	4	1	1	female	1	0
4	5	0	3	male	0	0


```
        = 5))
model.add(Dense(output_dim = 10, init = 'uniform', activation='relu'))
model.add(Dense(output_dim = 1, init= 'uniform', activation = 'sigmoid'))

model.compile(optimizer = 'adam', loss='binary_crossentropy',
              metrics=['accuracy'])
```

```
↳ /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:4: UserWarning: Update y
   after removing the cwd from sys.path.
   /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:5: UserWarning: Update y
   """
   /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:6: UserWarning: Update y
```

```
h = model.fit(X, y, epochs=100, batch_size=16, validation_split=0.2)
```

```
↳
```

Train on 712 samples, validate on 179 samples

Epoch 1/100

712/712 [=====] - 0s 582us/step - loss: 0.6910 - acc: 0.6138

Epoch 2/100

712/712 [=====] - 0s 95us/step - loss: 0.6795 - acc: 0.6096

Epoch 3/100

712/712 [=====] - 0s 91us/step - loss: 0.6465 - acc: 0.6629

Epoch 4/100

712/712 [=====] - 0s 101us/step - loss: 0.5935 - acc: 0.7837

Epoch 5/100

712/712 [=====] - 0s 93us/step - loss: 0.5437 - acc: 0.7879

Epoch 6/100

712/712 [=====] - 0s 93us/step - loss: 0.5122 - acc: 0.7963

Epoch 7/100

712/712 [=====] - 0s 75us/step - loss: 0.4921 - acc: 0.7921

Epoch 8/100

712/712 [=====] - 0s 73us/step - loss: 0.4804 - acc: 0.7949

Epoch 9/100

712/712 [=====] - 0s 74us/step - loss: 0.4722 - acc: 0.7878


h.history.keys()

dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])

Epoch 10/100

```
import matplotlib.pyplot as plt
plt.plot(h.history['loss'])
plt.plot(h.history['val_loss'])
```

[<matplotlib.lines.Line2D at 0x7ff9c3273fd0>]



Epoch 25/100

712/712 [=====]

0s 90us/step - loss: 0.4460 - acc: 0.8220

0.8220

Epoch 50/100

712/712 [=====] - 0s 89us/step - loss: 0.4460 - acc: 0.8220

0.8220

Epoch: 25/100

▼ MNIST

See how well a fully connected Neural Network performs on MNIST

```
from tensorflow.keras import datasets
(x_train, y_train), (x_test, y_test) = datasets.mnist.load_data()
```

↳ Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist11493376/11490434> [=====] - 0s 0us/step

```
import pandas as pd
```

```
x_train.shape
```

↳ (60000, 28, 28)

```
x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
y_train = pd.get_dummies(y_train)
y_test = pd.get_dummies(y_test)
```

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

```
model = Sequential()
```

```
model.add(Dense(20, input_shape=(784,), activation='relu'))
model.add(Dense(20, input_shape=(784,), activation='relu'))
model.add(Dense(10, activation='softmax'))
```

```
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

```
h = model.fit(x_train, y_train, epochs=100, batch_size=16, validation_split=0.2)
```

↳

Train on 48000 samples, validate on 12000 samples

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/

Instructions for updating:

Use tf.cast instead.

Epoch 1/100

48000/48000 [=====] - 6s 120us/sample - loss: 2.2455 - acc:

Epoch 2/100

48000/48000 [=====] - 5s 112us/sample - loss: 0.9709 - acc:

Epoch 3/100

48000/48000 [=====] - 5s 111us/sample - loss: 0.8360 - acc:

Epoch 4/100

48000/48000 [=====] - 5s 111us/sample - loss: 0.7711 - acc:

Epoch 5/100

48000/48000 [=====] - 5s 112us/sample - loss: 0.7123 - acc:

Epoch 6/100

48000/48000 [=====] - 5s 112us/sample - loss: 0.6707 - acc:

Epoch 7/100

48000/48000 [=====] - 5s 112us/sample - loss: 0.6225 - acc:

Epoch 8/100

48000/48000 [=====] - 5s 111us/sample - loss: 0.5688 - acc:

Epoch 9/100

48000/48000 [=====] - 5s 111us/sample - loss: 0.5145 - acc:

Epoch 10/100

48000/48000 [=====] - 6s 115us/sample - loss: 0.3581 - acc:

Epoch 11/100

48000/48000 [=====] - 5s 112us/sample - loss: 0.3103 - acc:

Epoch 12/100

48000/48000 [=====] - 5s 113us/sample - loss: 0.2847 - acc:

Epoch 13/100

48000/48000 [=====] - 5s 112us/sample - loss: 0.2735 - acc:

Epoch 14/100

48000/48000 [=====] - 5s 111us/sample - loss: 0.2634 - acc:

Epoch 15/100

48000/48000 [=====] - 5s 113us/sample - loss: 0.2541 - acc:

Epoch 16/100

48000/48000 [=====] - 5s 112us/sample - loss: 0.2478 - acc:

Epoch 17/100

48000/48000 [=====] - 5s 111us/sample - loss: 0.2393 - acc:

Epoch 18/100

48000/48000 [=====] - 5s 112us/sample - loss: 0.2351 - acc:

Epoch 19/100

48000/48000 [=====] - 5s 113us/sample - loss: 0.2253 - acc:

Epoch 20/100

48000/48000 [=====] - 5s 111us/sample - loss: 0.2278 - acc:

Epoch 21/100

48000/48000 [=====] - 5s 114us/sample - loss: 0.2185 - acc:

Epoch 22/100

48000/48000 [=====] - 5s 113us/sample - loss: 0.2154 - acc:

Epoch 23/100

48000/48000 [=====] - 5s 112us/sample - loss: 0.2069 - acc:

```
import matplotlib.pyplot as plt
plt.plot(h.history['loss'])
plt.plot(h.history['val_loss'])
plt.plot(h.history['acc'])
```



```
[<matplotlib.lines.Line2D at 0x7f44b219a048>]
```

